

Disciplina: Blockchain for Education

Kirill Kuvshinov¹, Jonn Mostovoy², and Ilya Nikiforov¹

¹Teach Me Please, <https://teachmeplease.com>

²Serokell, <https://serokell.io>

Version 0.4.1
December 1, 2017

Abstract

In this paper we analyze the main issues that arise from storing educational records in blockchain and propose the architecture of the Disciplina platform – a domain-specific blockchain implementation. The platform is designed to act as a decentralized ledger, with special regard for privacy and mechanisms of data disclosure. We present an overview of the main entities, their roles and incentives to support the network. Please note that the project is a work-in-progress and the descriptions provided are subject to change.

1 Introduction

Recent advances in blockchain technology and decentralized consensus systems open up new possibilities for building untamperable domain-specific ledgers with no central authority. Since the launch of Bitcoin [10] blockchains had been primarily used as a mechanism for value transfers. With the growth of the Ethereum platform [13], the community realized that by using a chain of blocks and consensus rules one can not only store value and track its movement, but, more generally, store some state and enforce conditions upon which this state can be modified.

Bitcoin, Ethereum and other permissionless blockchains were developed with the assumption that everyone is free to join the network and validate transactions, that are public. However, the industry often requires privacy, and thus the permissive solutions with private ledgers came to exist. These solutions include Tendermint [7], Hyperledger [4], Kadena [1] and others.

The increased interest and the variety of the blockchain technologies lead to the growth of their application domains. The idea of storing educational records in the blockchain has been circulating in the press and academic papers for several years. For example, [12] and [5] focus on the online education and propose to create a system based on the educational smart contracts in a public ledger. Recently, Sony announced a project that aims at incorporating educational records in a permissioned blockchain based on Hyperledger [2]. The ledger is going to be shared between major offline educational institutes.

The main issue these solutions have in common is that they target a certain subset of ways people get knowledge. We propose a more general approach that would unite the records of large universities, small institutes, schools and online educational platforms to form a publicly verifiable chain. Contrary to the solutions like Ethereum, we do not aim at proposing a programmable blockchain that fits all the possible applications. Rather, we believe, that we should harness all the latest knowledge that emerged in the last few years in the fields of consensus protocols, authenticated data structures and distributed computations to offer a new domain-specific ledger. In this paper we introduce Disciplina — the platform based on blockchain technology that aims to transform the way educational records are generated, stored and accessed.

2 Architecture overview

Due to the nature of the platform, it has to operate on sensitive data, such as courses, assignments, solutions and scores. Permissionless blockchains, like Ethereum or EOS, would require disclosing this data to the public, whereas the permissive ones, like Hyperledger, lack public verifiability. Our

architecture splits the blockchain into two layers: the private layer contains sensitive data, and the public one contains the information necessary to validate the integrity and authenticity of the private blocks. The key entities of the proposed blockchain architecture are presented in Figure 1.

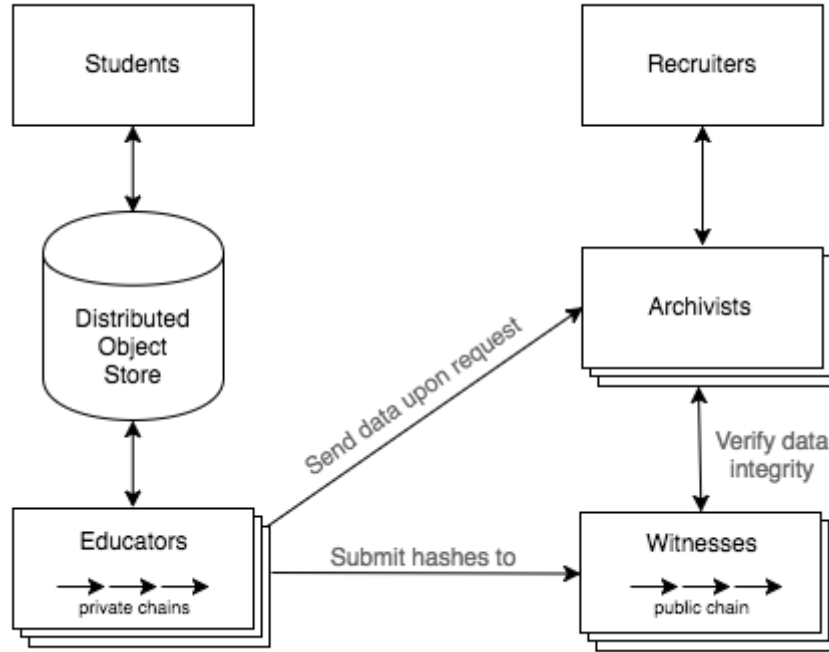


Figure 1: Key entities of the Disciplina platform

The private layer is maintained by each Educator independently of others. Educators can be either large educational institutes, capable of running their own nodes, or some trusted party that runs the chain for the self-employed teachers and small institutions. This layer contains the personalized information on the interactions between the students and the Educator. All the interactions, such as receiving an assignment, submitting solutions, or being graded, are treated as transactions in the private chain.

Students get access to the platform through web and mobile applications. Using the applications they choose Educators, enroll in courses, get assignments and submit solutions. The scores and the criteria of whether the Student has finished the course successfully are determined by the Educator. The education process from the platform's perspective is as follows:

1. A Student chooses an Educator and a course that she wants to enroll in
2. If the course is offered on a pre-paid basis, the Student uses her app to pay the fee.
3. The Student's application communicates with the Educators' software and performs the key exchange.
4. During the course, the Educator provides assignments that the Student has to complete in order to get the score. The Educator's software encrypts the assignments with the key the parties agreed upon, and puts it into a distributed Object Store.
5. The Student acquires the assignment, completes it and puts the solution to the Object Store (which is encrypted as well).
6. The Educator then grades the solution and transfers the artifacts along with the score to the blockchain.
7. Upon the completion of the course, the Student acquires a final score based on the scores she got for her assignments. This final score is also added to the Educator's chain.

Making the Educators’ chains private opens the possibility for Educators to tamper with the data in their chains. To overcome this issue and make the private transactions publicly verifiable, we introduce the second, public, layer of the blockchain. The public part of the network consists of Witnesses – the special entities that witness the fact that a private block was produced by Educator. They do so by writing the authentication information of a private block into the public chain, which is used in the future by an arbitrary Verifier to substantiate a proof of transaction inclusion given to it by an Archivist or an Educator. Witnesses also process public information issued by the Educators, such as announcements that an Educator has started or stopped offering a course in a particular discipline. The Witnesses agree on which public blocks are valid using the specified consensus rules.

Archivists are entities that provide an interface for the Recruiters to communicate with the platform. The Archivists act like a bridge between the Recruiters and the Educators: they choose the institutes that are relevant for the particular request, orchestrate the data disclosure and provide the evidence on the validity of that data. They obtain this evidence by communicating with Witnesses and comparing the headers of the private data blocks with the headers stored in the public Witnesses’ chain.

3 Implementation choices

In this section we describe the proposed architecture in more detail. We present the excerpt on the internal structure of both public and private chains and the reasoning behind these choices.

In order to deduce the internal structure of our system, we will first analyze its use-cases. The overview of the education process is given in section 2. The communication between the Student and the Educator is saved as transactions in the private chain. However, the implementation details of this chain mostly depend on the data disclosure process.

We will start from analyzing this process and determining the main issues that arise from the need to disclose and verify the validity of the private blocks. Then we will propose the structure of the private and public blocks that addresses these issues.

3.1 Anonymity and certification

The permissionless nature of our public chain leads to the ability for malevolent students to create educational institutes in order to get the scores for the courses they did not attend. Moreover, the knowledge students actually get by completing the course, and the conditions upon which the course is considered completed, vary significantly between the educational institutions.

These issues currently can not be solved solely on the protocol level: they require an external source of information to determine the physical existence and the reputation of an Educator. Although we leave the public chain open for the Educators to submit their private block headers, we propose to add a separate layer of reputation and trust on top of the protocol.

We do so by introducing the Archivists – the entities that join the network with the approval from another Archivist. Their main role is to add a trust level above just the raw protocol: they store the certificates of the Educational institutes and have the right to revoke those certificates. Furthermore, the Archivists are the entities responsible for determining the ratings of the particular Educators. The Archivists base their rating on the off-chain sources of information and gain authority for providing valid ratings and performing all the necessary compliance procedures for the Educators.

While in theory the Recruiters can query the Educators directly and initiate a data disclosure request, they are generally discouraged to do so. We expect the Recruiters to be willing to pay extra fees to the Archivists for certificate validation and ranking of the Students depending on the Educators’ ratings and other factors upon request. We describe the data disclosure process in detail in section 3.7

3.2 Activity Type Graph

When a Recruiter makes a request to one of the Archivists, the Archivist has to somehow choose the relevant Educators. Moreover, when an Educator discloses the data, it has to provide as minimal set of entries as possible. This set has to be verifiable, which means that the Educator provides the proof of the data validity along with the data being disclosed.

In order to achieve these goals, we divide the data that the Educators store into atomic Activity Types. Each Educator maintains a journal of transactions per each Activity Type that the Educator offers.

All the Activity Types are grouped into courses that are further grouped into larger entities such as subjects and areas of knowledge. This grouping can be stored as the Activity Type Graph G_A with the following properties:

1° G_A is a directed graph:

$$G_A : \langle V : \{\text{Vert}\}, e_{out} : \text{Vert} \rightarrow \{\text{Vert}\} \mid \text{rest} \rangle \quad (1)$$

2° Each vertex of G_A is associated with depth:

$$G_A : \langle d : \text{Vert} \rightarrow \text{Int} \mid \text{rest} \rangle \quad (2)$$

3° Law of pointing down:

$$G_A : \langle v \in e_{out}(u) \implies d(v) > d(u) \rangle \quad (3)$$

4° G_A has special *et cetera* vertices u :

$$\forall v \in V \exists u (u \in e_{out}(v) \wedge e_{out}(u) = \emptyset) \quad (4)$$

The example of the Activity Type Graph (ATG) is shown in Figure 2. The vertex v of the graph is a *leaf* if $e_{out}(v) = \emptyset$. Otherwise we call it an *internal vertex*. Every internal vertex of the graph has a special *etc.* child (some of these are omitted in the figure).

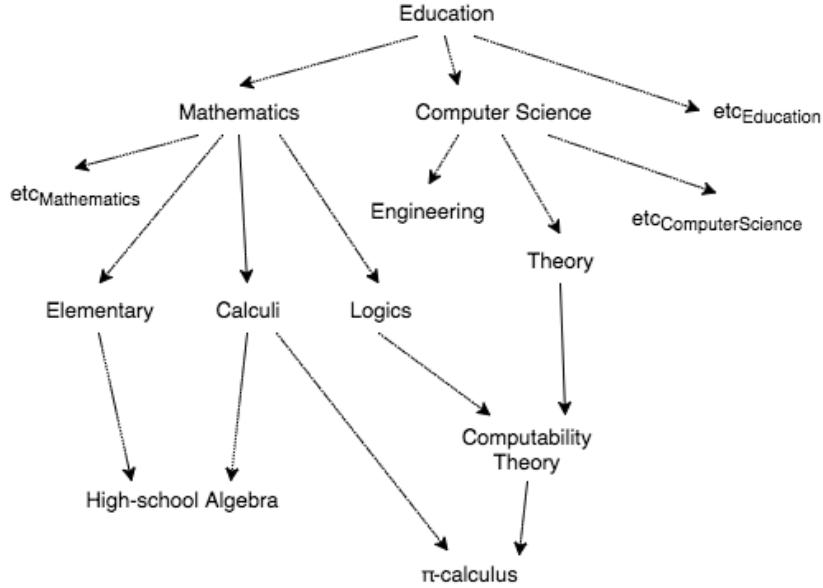


Figure 2: An example of the Activity Type Graph. Some of the vertices are not shown

The need for *etc.* vertices arises from the fact that not all of the Educators teach courses exactly in leaves — some of them offer general courses that provide just the necessary background. For example, some of the universities teach the basic “Computer science” course, that contains the basics of the discipline. In this case, when the particular category is hard to define, the university would use the `etcComputerScience` vertex.

On the protocol level, the Educators can announce that they teach a particular course, but can not modify the Activity Type Graph structure. The structure of the graph is maintained by the core developers and updated upon request from the Educators.

3.3 Private chain

The structure of the private block is shown in Figure 3. The block consists of a public *header* that the Educators relay to the Witnesses, and the private *body* that remains in the educational institute until it receives a data disclosure request.

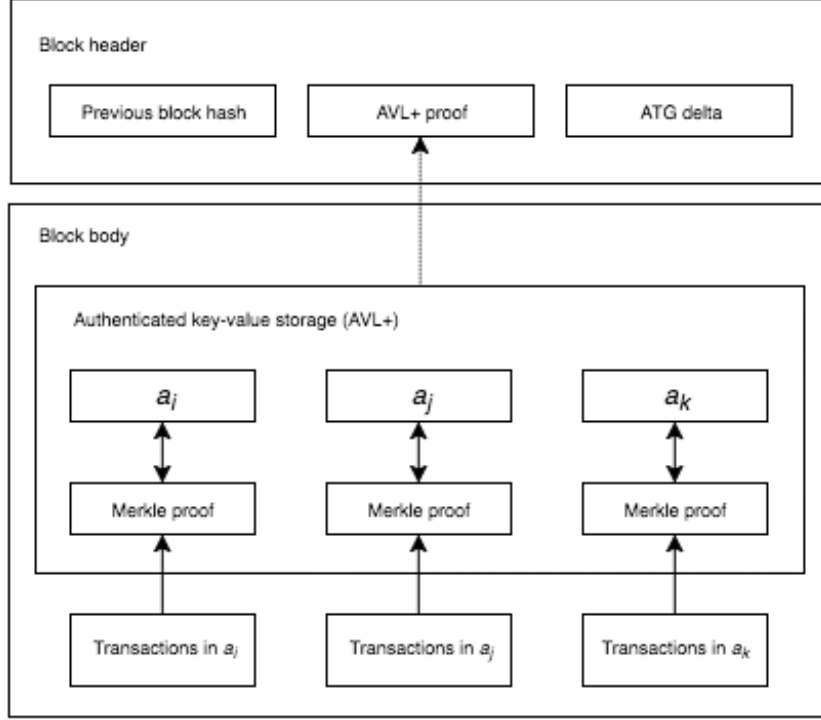


Figure 3: Private block structure

During the educational process the Educators emit atomic *private transactions*. These transactions represent the modifications to the journal of academic achievements in a certain activity type (thus, making a transaction means appending the data to the journal of a particular course). The transactions can be of the following types:

- student enrolls in a course;
- student gets an assignment;
- student submits an assignment;
- student gets a score for an assignment;
- student gets a final score for the course.

Let us denote an i -th transaction belonging to a certain activity type a_j as $T_{a_j}^i$. The Educators group the transactions that occurred during the current block time slot according to the activity type, and construct Merkle trees [9] for these journal modifications:

$$M_{a_j} = \text{MerkleTree}(T_{a_j}^i) \quad (5)$$

The Educator's private block body contains a dictionary of key-value pairs, where each key is the leaf in the activity type graph a_j , and the values are the Merkle roots of all the transactions that occurred in this leaf in the current block. Thus, the private block body is a mapping:

$$a_j \rightarrow \text{root}(M_{a_j}) \quad (6)$$

In order to make this mapping easily verifiable, we use a structure called the *authenticated AVL+ tree* introduced in [11]. This structure is based on state-of-the-art research that enables for

faster verification of the mapping and allows us to never disclose it: the Witnesses would not have to store the whole blockchain like Bitcoin or Ethereum nodes do. Rather, they would just have to check the private block headers in order to confirm that none of the private blocks were tampered with.

The private block header consists of the AVL+ proof along with the previous block hash and the information on the Activity Type Graph modifications (ATG delta). The *ATG delta* part allows the Educators to inform the Witnesses and the Archivists of the modifications to the courses they teach.

After the end of the block time slot, an Educator signs the block header and submits it to the Witnesses so that the private transactions can be confirmed by the public chain. Thus, the private blocks form a publicly verifiable chain of events, grouped according to the activity type.

3.4 Public chain

The Witnesses maintain a public chain – a distributed ledger that contains publicly available information. If one wishes to perform a transaction on the public chain, she has to pay a certain fee that serves two purposes. First of all, the fee incentivizes the Witnesses to participate in the network and issue new blocks. Second, by requiring a fee for each transaction, we protect the public ledger from being spammed.

We present the structure of the public blocks in Figure 4. The public ledger contains the following information:

1. Modification history of the Activity Type Graph.
2. Private transaction proofs.
3. Account balances and value transfer history.

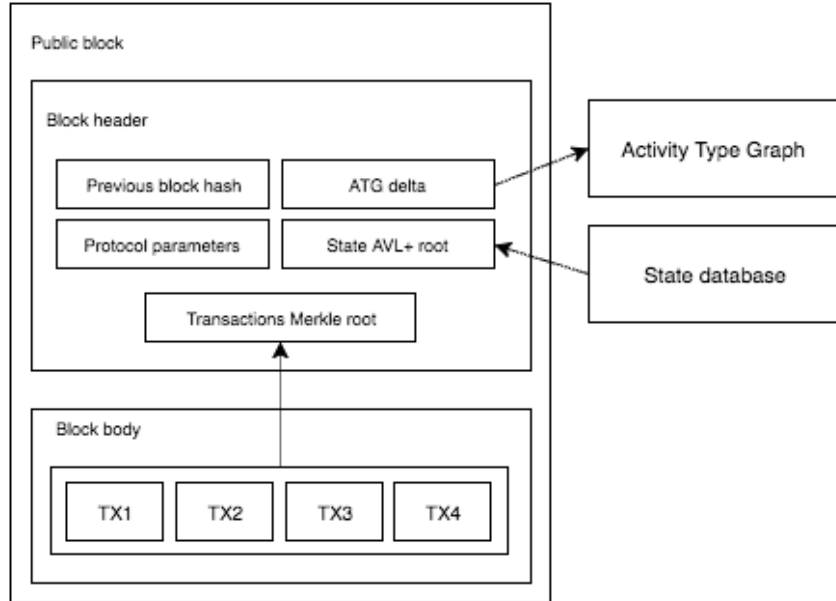


Figure 4: Public block structure

There are two major ways to store the account balances and other state information: UTXO and account-based architectures. UTXO is an unspent transaction output, that contains some predicate – a condition that has to be fulfilled in order to spend the coins. To prove the money ownership, the spender provides a witness – an input that makes the predicate true. Thus, the UTXO-based architecture requires the transactions to be stateless, effectively limiting the application domain [3]. The unspent outputs with an associated state can be treated as smart-contracts in the account-based architectures like Ethereum [13]. The state is stored in an off-chain storage – the state

database. The transactions are treated as the modifications of the world state. We are going to develop an UTXO-based ledger with transactions programmable in Plutus language [8].

The recent achievements in the field of consensus protocols, like the provably secure Ouroboros [6], allow us to build a public chain based on the Proof of Stake consensus rules. Thus, we can increase the transaction speed and drop the need for the expensive mining. However, with mining being dropped, we need to provide incentives for the Witnesses to maintain the chain and participate in the network.

3.5 Incentives

In order to incentivize the Witnesses as well as the Archivists and the Object Store maintainers we propose a monetary policy with two main sources of income. The first one is the technical pool — a special pool of tokens, which are reserved until the participants acquire them through contributions to the operation of the platform. The tokens from the technical pool will be distributed with exponential slowdown. Running the nodes for different entities of the system require different hardware resources and there may come a point where the system lacks the nodes of a certain entity. To overcome this issue, the complexity and the amount of tokens received by the participants will be determined dynamically so that the equilibrium between the entities is preserved, for example, if the system lacks the Archivists, the incentive to run the Archivist's node would be more then the one for the Witnesses.

The second source of the participants' income is the fees for the transactions in the system. The Recruiters' fee is distributed among all the participants except the Object Store maintainers. The latter obtain tokens from the Educators paying them for the storage they offer.

3.6 Fair CV

One of the main goals of the Disciplina platform is to provide a way for the Students to easily prove their educational records. We propose to duplicate the records in the Student's *digital CV*. This CV contains all the records that the parties have generated during the Student's educational process along with the validity proofs of that data (see figure 5).

In order to prove that the transaction actually occurred in some private block of a particular Educator, the student has to provide the cryptographic proofs along with the actual data. The cryptographic proof of the inclusion of an element in an authenticated data structure is generally a path of hashes. Let us denote the path of the element e in the data structure D as $\text{path}(e, D)$. Thus, the Student has to provide the following data:

- The Student's and the Educator's public keys PK_S and PK_E .
- The course a_i and the a private transaction T_{priv} with the score.
- The Merkle path of the transaction in the journal: $P_M = \text{path}(T_{priv}, M_{a_i})$
- The AVL+ path to prove that the journal was included in a private block's AVL+ tree A : $P_A = \text{path}(\text{root}(M_{a_i}), A)$
- The public block number H and the Merkle path of the transaction T_{pub} that pushed the private block into the public chain: $P_H = \text{path}(T_{pub}, M_H)$, where M_H is a Merkle tree of the transactions in the block H .

Having this data one can prove the occurrence of a certain transaction in one of the Educator's private blocks without the need to request any data from the Educator during the validation process. Thus, any party can check the validity of the Student's CV for free if the Student wishes to disclose it.

Let $\rho(e, P)$ be the function that substitutes the element e in path P and computes the root hash of the authenticated data structure. Then the validation process is as follows:

1. Query the public chain to find the block H and obtain the Merkle root of the transactions: $\text{root}(M_H)$.
2. Check whether $\rho(T_{pub}, P_H) = \text{root}(M_H)$.
3. Check that the public transaction T_{pub} was signed with the Educator's public key PK_E .

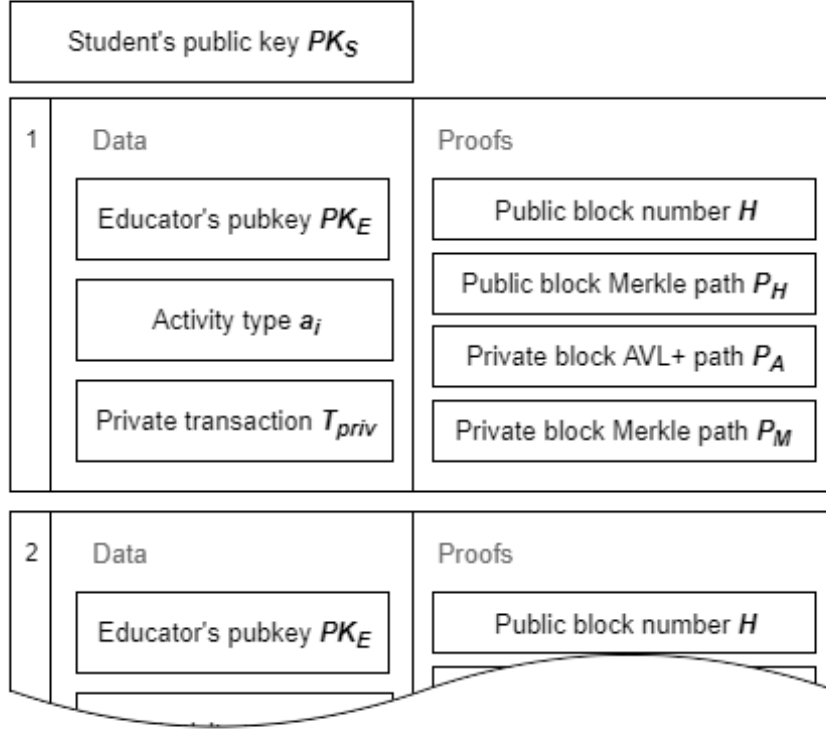


Figure 5: Student's authenticated CV

4. From the public transaction T_{pub} obtain the root of the AVL+ tree: $\text{root}(A)$.
5. Use PK_S , a_j and score to build a private transaction T_{priv} .
6. Compute $R_M = \rho(T_{priv}, P_M)$.
7. Check that $\rho(R_M, P_A) = \text{root}(A)$.

These validation steps can prove that an Educator with a public key PK_E issued a transaction T_{priv} in one of its private blocks. One can attribute the PK_E to a particular real-world educational institution by checking the Educator's certificate as described in section 3.1.

3.7 Data Disclosure

The high-level data disclosure process is as follows:

1. Archivists keep records on all the Educators in the system along with the courses they teach.
2. The Recruiter sends data disclosure request to one of the Archivists.
3. The Archivist selects the relevant Educators and forwards the request.
4. The Educators determines the size of the data they are going to disclose.
5. Based on the size of the data, the Archivist informs the Recruiter of the cost of the request.
6. The Recruiter pays the fee.
7. The Educators send the data to the Archivist.
8. The Archivist checks the blocks received from the Educator for their validity using the data from the Witnesses' chain, and forwards this data to the Recruiter.

9. The fee is distributed between the relevant parties.

The process of data disclosure involves direct communication between a particular Educator, willing to disclose a part of the data, and a Recruiter, willing to pay for this data. To mitigate the risk of secondary market creation, one should ensure that the majority of the data remains in the private blocks. Thus we reduce the size of the Educators' response by incentivizing the Recruiters to make as accurate requests as possible. We propose the following formula to determine the cost of the request:

$$\text{Cost}(\text{request}) \sim \exp(N_{\text{actions}}(\text{response})) \quad (7)$$

The value received from the Recruiter is then distributed between the Archivists, the Educators that disclosed the data, and the affected students.

4 Conclusion

In this paper we presented the architecture of the Disciplina platform. The described architecture provides a way to store educational records in the blockchain while preserving the privacy of these records. The concepts of private chains and a digital CV make it possible to verify the educational records of a particular person. The Archivist entities provide reputational semantics to the educational institutes listed in the digital CVs.

We developed our platform not only as the source of trust, but also as a database of the students from all over the world. We believe that the data that is stored in the system has a value in itself. The need to disclose this data was also addressed in the paper: we described a mechanism for the fair data trade and the measures against the secondary market creation.

References

- [1] Kadena: Scalable blockchain. <https://kadena.io/>.
- [2] Sony global education. <https://www.sonyged.com/>.
- [3] I. Bentov, R. Kumaresan, and A. Miller. Instantaneous decentralized poker. *arXiv preprint arXiv:1701.06726*, 2017.
- [4] C. Cachin. Architecture of the hyperledger blockchain fabric. In *Workshop on Distributed Cryptocurrencies and Consensus Ledgers*, 2016.
- [5] P. Devine. Blockchain learning: can crypto-currency methods be appropriated to enhance online learning? 2015.
- [6] A. Kiayias, A. Russell, B. David, and R. Oliynykov. Ouroboros: A provably secure proof-of-stake blockchain protocol. In *Annual International Cryptology Conference*, pages 357–388. Springer, 2017.
- [7] J. Kwon. Tendermint: Consensus without mining. *Draft v. 0.6, fall*, 2014.
- [8] D. McAdams. Formal specification of the plutus language. <https://github.com/input-output-hk/plutus-prototype/blob/master/docs/spec/Formal%20Specification%20of%20the%20Plutus%20Language%20-%20McAdams.pdf>.
- [9] R. C. Merkle. A certified digital signature. In *Conference on the Theory and Application of Cryptology*, pages 218–238. Springer, 1989.
- [10] S. Nakamoto. Bitcoin: A peer-to-peer electronic cash system, 2008.
- [11] L. Reyzin, D. Meshkov, A. Chepuronov, and S. Ivanov. Improving authenticated dynamic dictionaries, with applications to cryptocurrencies. *IACR Cryptology ePrint Archive*, 2016:994, 2016.
- [12] M. Swan. *Blockchain: Blueprint for a new economy*. O'Reilly Media, Inc., 2015.
- [13] G. Wood. Ethereum: A secure decentralised generalised transaction ledger. *Ethereum Project Yellow Paper*, 151, 2014.